

```
In [2]: import pandas as pd
import numpy as np

raw = pd.read_csv("boston.txt", sep=r"\s+", skiprows=22, header=None)
X = np.hstack([raw.values[:, 2:], raw.values[1:, :2]])
y = raw.values[1:, 2]
columns = ["crim", "zn", "indus", "chas", "nox", "rm", "age",
           "dis", "rad", "tax", "ptratio", "b", "lstat"]
df = pd.DataFrame(X, columns=columns)
df["medv"] = y
```

a

```
In [5]: import matplotlib.pyplot as plt
from seaborn import pairplot
import statsmodels.api as sm

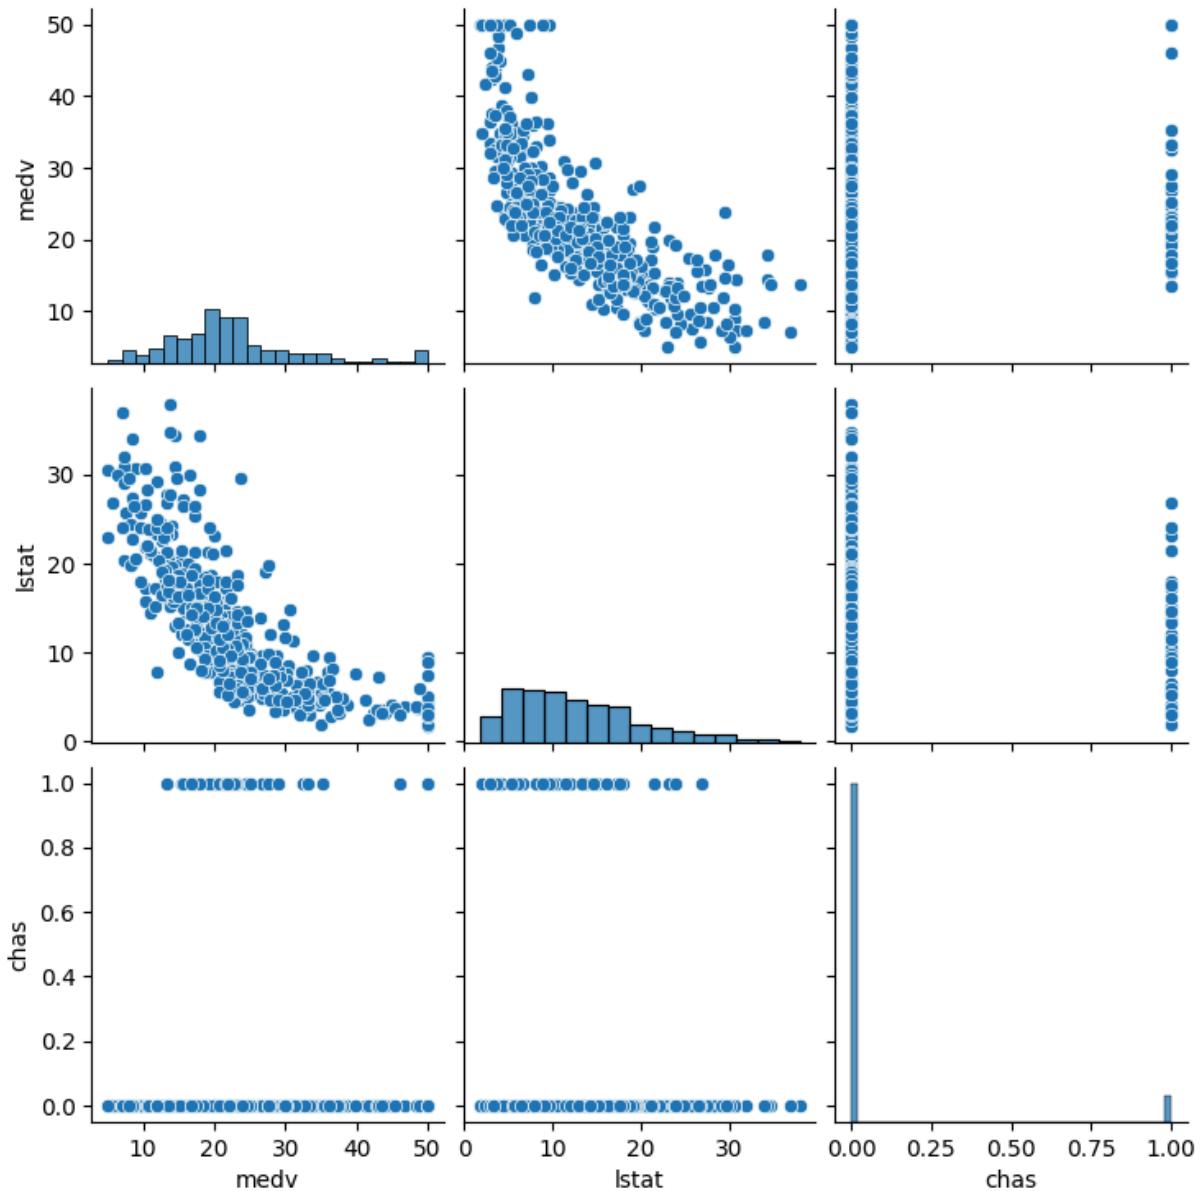
pairplot(df[["medv", "lstat", "chas"]])
X = sm.add_constant(df[["lstat", "chas"]])
model = sm.OLS(df["medv"], X).fit()
print(model.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.563
Model:	OLS	Adj. R-squared:	0.561
Method:	Least Squares	F-statistic:	323.4
Date:	Tue, 10 Feb 2026	Prob (F-statistic):	4.93e-91
Time:	17:06:14	Log-Likelihood:	-1631.1
No. Observations:	506	AIC:	3268.
Df Residuals:	503	BIC:	3281.
Df Model:	2		
Covariance Type:	nonrobust		
<hr/>			
	coef	std err	t
const	34.0941	0.561	60.809
lstat	-0.9406	0.038	-24.729
chas	4.9200	1.069	4.601
<hr/>			
Omnibus:	131.896	Durbin-Watson:	0.975
Prob(Omnibus):	0.000	Jarque-Bera (JB):	275.510
Skew:	1.406	Prob(JB):	1.49e-60
Kurtosis:	5.272	Cond. No.	57.8
<hr/>			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



The pair plots indicate a non linear relationship between `lstat` and `medv`, `medv` drops quickly for low-mid `lstat`, and then flattens out. Usually taking the log, square root, or inverse of `lstat` can help fix this. For this problem, we will use a log transformation.

Intercept: 34.0941, SE 0.561, p < 0.0001

This is the predicted `medv` for an `lstat` and `chas` value of 0. This isn't meaningful realistically as `lstat` 0 isn't possible, but this anchors the regression line. The value is statistically significant and well estimated from the low SE and p value.

`lstat` coefficient: -0.9406, SE 0.038, p < 0.0001

Holding `chas` constant, a 1 percentage point increase in `lstat` is attributed to about a \$940.6 decrease in `medv`. The value is statistically significant and well estimated from the low SE and p value.

`chas` coefficient: 4.92, SE 1.069, p < 0.0001 Holding `lstat` constant, properties that border the Charles River are attributed to a \$4,920 increase in `medv`. The value is

statistically significant and well estimated from the low SE and p value.

b

In [6]: `# add interaction term to model`

```
df["lstat_chas"] = df["lstat"] * df["chas"]
X = sm.add_constant(df[["lstat", "chas", "lstat_chas"]])
model_interaction = sm.OLS(df["medv"], X).fit()
print(model_interaction.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.569			
Model:	OLS	Adj. R-squared:	0.566			
Method:	Least Squares	F-statistic:	220.8			
Date:	Tue, 10 Feb 2026	Prob (F-statistic):	2.68e-91			
Time:	17:16:58	Log-Likelihood:	-1627.4			
No. Observations:	506	AIC:	3263.			
Df Residuals:	502	BIC:	3280.			
Df Model:	3					
Covariance Type:	nonrobust					
<hr/>						
	coef	std err	t	P> t	[0.025	0.975]
const	33.7672	0.570	59.222	0.000	32.647	34.887
lstat	-0.9150	0.039	-23.478	0.000	-0.992	-0.838
chas	9.8251	2.103	4.672	0.000	5.693	13.957
lstat_chas	-0.4329	0.160	-2.703	0.007	-0.748	-0.118
<hr/>						
Omnibus:	130.570	Durbin-Watson:	1.007			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	275.834			
Skew:	1.383	Prob(JB):	1.27e-60			
Kurtosis:	5.330	Cond. No.	114.			
<hr/>						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Intercept: 33.7672, SE 0.570, p < 0.0001

This is the predicted `medv` for an `lstat` and `chas` value of 0. This isn't meaningful realistically as `lstat` 0 isn't possible, but this anchors the regression line. The value is statistically significant and well estimated from the low SE and p value.

lstat coefficient: -0.9150, SE 0.039, p < 0.0001

Holding `chas` constant, a 1 percentage point increase in `lstat` is attributed to about a \$915 decrease in `medv`. The value is statistically significant and well estimated from the low SE and p value. The coefficient is slightly smaller than the model without the interaction term.

chas coefficient: 9.8251, SE 2.103, p < 0.0001 Holding `lstat` constant, properties that border the Charles River are attributed to a \$9,825.10 increase in `medv`. The value is statistically significant and well estimated from the low SE and p value. The coefficient is significantly larger than the model without the interaction term.

lstat:chas coefficient: -0.4329, SE 0.160, p 0.007

The negative effect of `lstat` is stronger for properties bordering the Charles River. The value is statistically significant and well estimated from the low SE and p value. Specifically, for properties that border the Charles River, a 1 percentage point increase in `lstat` decreased `medv` by 1,347.90 compared to the 915 elsewhere.

Given that the new R squared value of this model (0.569, 0.566) is barely an improvement from the previous model (0.563, 0.561), interactions do not improve the model, but the term is real and statistically meaningful.

C

```
In [7]: # fit two separate models of medv as a function of lstat
# one model for chas = 0 and one model for chas = 1
# compare to joint interaction model to determine which is better

df0 = df[df["chas"] == 0]
df1 = df[df["chas"] == 1]

X0 = sm.add_constant(df0["lstat"])
model0 = sm.OLS(df0["medv"], X0).fit()
print(model0.summary())

X1 = sm.add_constant(df1["lstat"])
model1 = sm.OLS(df1["medv"], X1).fit()
print(model1.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.552
Model:	OLS	Adj. R-squared:	0.551
Method:	Least Squares	F-statistic:	577.2
Date:	Tue, 10 Feb 2026	Prob (F-statistic):	9.59e-84
Time:	17:25:11	Log-Likelihood:	-1504.9
No. Observations:	471	AIC:	3014.
Df Residuals:	469	BIC:	3022.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	33.7672	0.557	60.603	0.000	32.672	34.862
lstat	-0.9150	0.038	-24.026	0.000	-0.990	-0.840

Omnibus:	139.989	Durbin-Watson:	1.006
Prob(Omnibus):	0.000	Jarque-Bera (JB):	330.529
Skew:	1.523	Prob(JB):	1.68e-72
Kurtosis:	5.751	Cond. No.	30.0

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.582
Model:	OLS	Adj. R-squared:	0.569
Method:	Least Squares	F-statistic:	45.90
Date:	Tue, 10 Feb 2026	Prob (F-statistic):	1.01e-07
Time:	17:25:11	Log-Likelihood:	-120.33
No. Observations:	35	AIC:	244.7
Df Residuals:	33	BIC:	247.8
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	43.5923	2.593	16.814	0.000	38.318	48.867
lstat	-1.3479	0.199	-6.775	0.000	-1.753	-0.943

Omnibus:	1.398	Durbin-Watson:	1.269
Prob(Omnibus):	0.497	Jarque-Bera (JB):	1.349
Skew:	0.395	Prob(JB):	0.509
Kurtosis:	2.453	Cond. No.	25.9

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Model	n	R^2	Adj R^2	AIC	BIC
Interaction	506	0.569	0.566	3263	3281

Model	n	R^2	Adj R^2	AIC	BIC
chas = 0	471	0.552	0.551	3014	3022
chas = 1	35	0.582	0.569	244.7	247.8

Note that the slopes for each `chas` specific model are the same as the indicated slopes in the interaction model when accounting for the interaction term. Additionally, the `chas = 1` model only has 35 samples. Thus, we prefer the interaction model due to it using all 506 observations, and provides a direct statistical test for slope differences in the `chas` specific models with better estimates. The comparison given by the model is straightforward and captures the majority of the information given by both groups.

d

```
In [ ]: df_restrict = df[df["rad"] <= 8]

X_num = sm.add_constant(df_restrict[["lstat", "rad"]].astype(float))
model_num = sm.OLS(df_restrict["medv"].astype(float), X_num).fit()
print(model_num.summary())

rad_dummies = pd.get_dummies(df_restrict["rad"], prefix="rad")
X_cat = sm.add_constant(pd.concat([df_restrict["lstat"].astype(float), rad_dummies.
model_cat = sm.OLS(df_restrict["medv"].astype(float), X_cat).fit()
print(model_cat.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.458
Model:	OLS	Adj. R-squared:	0.455
Method:	Least Squares	F-statistic:	156.7
Date:	Tue, 10 Feb 2026	Prob (F-statistic):	4.77e-50
Time:	17:37:08	Log-Likelihood:	-1212.6
No. Observations:	374	AIC:	2431.
Df Residuals:	371	BIC:	2443.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	33.0491	1.096	30.164	0.000	30.895	35.204
lstat	-0.9524	0.054	-17.592	0.000	-1.059	-0.846
rad	0.3819	0.197	1.937	0.054	-0.006	0.770

Omnibus:	85.140	Durbin-Watson:	0.926
Prob(Omnibus):	0.000	Jarque-Bera (JB):	147.053
Skew:	1.313	Prob(JB):	1.17e-32
Kurtosis:	4.593	Cond. No.	44.0

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.500
Model:	OLS	Adj. R-squared:	0.489
Method:	Least Squares	F-statistic:	45.67
Date:	Tue, 10 Feb 2026	Prob (F-statistic):	1.40e-50
Time:	17:37:08	Log-Likelihood:	-1197.4
No. Observations:	374	AIC:	2413.
Df Residuals:	365	BIC:	2448.
Df Model:	8		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	30.5985	0.585	52.271	0.000	29.447	31.750
lstat	-0.9116	0.054	-16.761	0.000	-1.019	-0.805
rad_1.0	0.4850	1.239	0.391	0.696	-1.952	2.922
rad_2.0	5.3733	1.141	4.708	0.000	3.129	7.618
rad_3.0	5.6042	0.930	6.026	0.000	3.775	7.433
rad_4.0	1.9095	0.646	2.958	0.003	0.640	3.179
rad_5.0	4.8236	0.617	7.824	0.000	3.611	6.036
rad_6.0	1.5964	1.116	1.430	0.154	-0.599	3.792
rad_7.0	3.7895	1.335	2.839	0.005	1.165	6.414
rad_8.0	7.0170	1.139	6.159	0.000	4.777	9.257

Omnibus:	93.551	Durbin-Watson:	1.000
Prob(Omnibus):	0.000	Jarque-Bera (JB):	181.753
Skew:	1.350	Prob(JB):	3.41e-40
Kurtosis:	5.093	Cond. No.	1.10e+17

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 4.55e-30. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Model	n	R^2	Adj R^2	AIC	BIC
Numerical rad	374	0.458	0.455	2431	2443
Categorical rad	374	0.500	0.489	2413	2448

Using `rad` as categorical significantly improves R^2 by 0.4. The AIC strongly favors the categorical specification, as the delta of 18 shows that allowing non-linear, level-specific effects of `rad` improves the fit. BIC slightly favors numerical specification however, as `rad` is heavily penalized with extra parameters as a categorical variable. Looking at the coefficients, the numerical `rad` value is barely statistically significant with a p-value of 0.054. However for the categorical model, different `rads` have varying statistical significance levels. This matches how `rad` is constructed realistically as it is an index, not a continuous economic quantity. Thus overall, we prefer the categorical specification due to higher adjusted R^2 and better representation of the predictor variables.

e

```
In [13]: b3 = model_cat.params["rad_3.0"]
b4 = model_cat.params["rad_4.0"]
diff = b3 - b4
print(f"Estimated difference in expected value of medv between houses w/ rad = 3 and rad = 4: {diff}")

cov = model_cat.cov_params()
var_diff = (
    cov.loc["rad_3.0", "rad_3.0"]
    + cov.loc["rad_4.0", "rad_4.0"]
    - 2 * cov.loc["rad_3.0", "rad_4.0"]
)
se_diff = np.sqrt(var_diff)
print(f"Estimated standard error of the difference: {se_diff}")
```

Estimated difference in expected value of medv between houses w/ rad = 3 and rad = 4: 3.6947186903988825

Estimated standard error of the difference: 1.1451240401723881