

hw3_project

January 29, 2026

```
[12]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import statsmodels.api as sm
from scipy.stats import norm
```

```
[4]: insurance = pd.read_csv("../datasets/insurance.csv")
insurance.head()
```

```
[4]:    age      sex      bmi  children smoker      region      charges
 0   19  female  27.900        0     yes  southwest  16884.92400
 1   18     male  33.770        1     no  southeast  1725.55230
 2   28     male  33.000        3     no  southeast  4449.46200
 3   33     male  22.705        0     no  northwest  21984.47061
 4   32     male  28.880        0     no  northwest  3866.85520
```

1 1

```
[8]: # continuous predictor variables only
X = insurance[['age', 'bmi', 'children']]
y = insurance['charges']

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
print(model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:      charges      R-squared:       0.120
Model:              OLS          Adj. R-squared:  0.118
Method:             Least Squares  F-statistic:    60.69
Date:      Thu, 29 Jan 2026  Prob (F-statistic): 8.80e-37
Time:      19:32:28          Log-Likelihood: -14392.
No. Observations:  1338         AIC:            2.879e+04
Df Residuals:     1334         BIC:            2.881e+04
Df Model:           3
Covariance Type:  nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
<hr/>						
const	-6916.2433	1757.480	-3.935	0.000	-1.04e+04	-3468.518
age	239.9945	22.289	10.767	0.000	196.269	283.720
bmi	332.0834	51.310	6.472	0.000	231.425	432.741
children	542.8647	258.241	2.102	0.036	36.261	1049.468
<hr/>						
Omnibus:		325.395	Durbin-Watson:			2.012
Prob(Omnibus):		0.000	Jarque-Bera (JB):			603.372
Skew:		1.520	Prob(JB):			9.54e-132
Kurtosis:		4.255	Cond. No.			290.
<hr/>						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Intercept: -6916.2433, SE 1757.480, p < 0.0001

This is the predicted charge for a newborn (age 0) with a bmi of 0, no children, female, non-smoker, from the northeast region. This has no meaning in it of itself, however given its extremely low p-value, is obviously extremely relevant to the model as it anchors the regression line.

Age: 239.9945, SE 22.289, p < 0.0001 Holding all other variables constant, each year a patient is older is associated with an increase of roughly \$240 in insurance charge. The SE is low and p-value is also extremely low, indicating a decently estimated and statistically significant predictor.

BMI: 332.0834, SE 51.31, p < 0.0001 Holding all other variables constant, a one unit increase in BMI is associated with an increase of roughly \$332 in insurance charge. The SE is high and p-value is extremely low, indicating an average estimate but still statistically significant predictor.

Children: 542.8647, SE 258.241, p 0.036 Holding all other variables constant, every additional child is associated with an increase of roughly \$542 in insurance charge. The SE is higher than both coefficient estimates for *age* and *bmi*, but still statistically significant at an alpha level of 0.05.

The adjusted R-squared value of the model is 0.118. This means that after accounting for model complexity, only about 11.8% of the variation in insurance charges is explained by age, BMI, and number of children. This indicates that most variation in charges is driven by non-continuous predictor variables, which Lantz does wish to show.

2 2

```
[11]: # recompute standard error and p values with sandwich formula
```

```
robust_cov = model.get_robustcov_results(cov_type='HC1')
print(robust_cov.summary())
```

OLS Regression Results

Dep. Variable:	charges	R-squared:	0.120
Model:	OLS	Adj. R-squared:	0.118

```

Method: Least Squares F-statistic: 60.13
Date: Thu, 29 Jan 2026 Prob (F-statistic): 1.85e-36
Time: 19:40:39 Log-Likelihood: -14392.
No. Observations: 1338 AIC: 2.879e+04
Df Residuals: 1334 BIC: 2.881e+04
Df Model: 3
Covariance Type: HC1
=====
            coef    std err      t    P>|t|    [0.025    0.975]
-----
const     -6916.2433  1703.922   -4.059    0.000  -1.03e+04  -3573.585
age        239.9945   22.355    10.736    0.000   196.140   283.849
bmi        332.0834   56.384     5.890    0.000   221.472   442.695
children   542.8647  243.126     2.233    0.026   65.913   1019.816
=====
Omnibus: 325.395 Durbin-Watson: 2.012
Prob(Omnibus): 0.000 Jarque-Bera (JB): 603.372
Skew: 1.520 Prob(JB): 9.54e-132
Kurtosis: 4.255 Cond. No. 290.
=====
```

Notes:

[1] Standard Errors are heteroscedasticity robust (HC1)

All coefficient values are the same, but this is expected as the sandwich formula does not change the fitted coefficients, only the estimated covariance matrix.

Intercept: The SE decreased slightly, however p-value still indicates a statistically significant value. This means that heteroskedasticity correction slightly tightens uncertainty around the intercept.

Age: Value remains relatively unchanged, thus the effect of age is very stable and not sensitive to heteroskedasticity.

BMI: The SE has increased, however p-value still indicates a statistically significant predictor variable. This means that OLS underestimated uncertainty for BMI. After correcting, BMI is still significant but less precisely estimated.

Children: The SE decreased slightly, and p-value has decreased as well. This means that once heteroskedasticity is accounted for, the effect of children is actually slightly more statistically significant.

Because all coefficient values are the same, the adjusted R-squared is still 0.118.

3 3

```
[13]: x = insurance['age']
y = insurance['charges']
conf = 0.95

x = np.asarray(x)
```

```

y = np.asarray(y)

X = sm.add_constant(x)
model = sm.OLS(y, X).fit()

xg = np.linspace(min(x), max(x), 100)
Xg = sm.add_constant(xg)
yhat_g = model.predict(Xg)

V_std = model.cov_params()
V_rob = model.get_robustcov_results(cov_type='HC1').cov_params()

se_std = np.sqrt(np.sum(Xg @ V_std * Xg, axis=1))
se_rob = np.sqrt(np.sum(Xg @ V_rob * Xg, axis=1))

z_point = norm.ppf(1 - (1 - conf) / 2)
z_bonf = norm.ppf(1 - (1 - conf) / (2 * len(xg)))

pw_std_u = yhat_g + z_point * se_std
pw_std_l = yhat_g - z_point * se_std

pw_rob_u = yhat_g + z_point * se_rob
pw_rob_l = yhat_g - z_point * se_rob

bonf_std_u = yhat_g + z_bonf * se_std
bonf_std_l = yhat_g - z_bonf * se_std

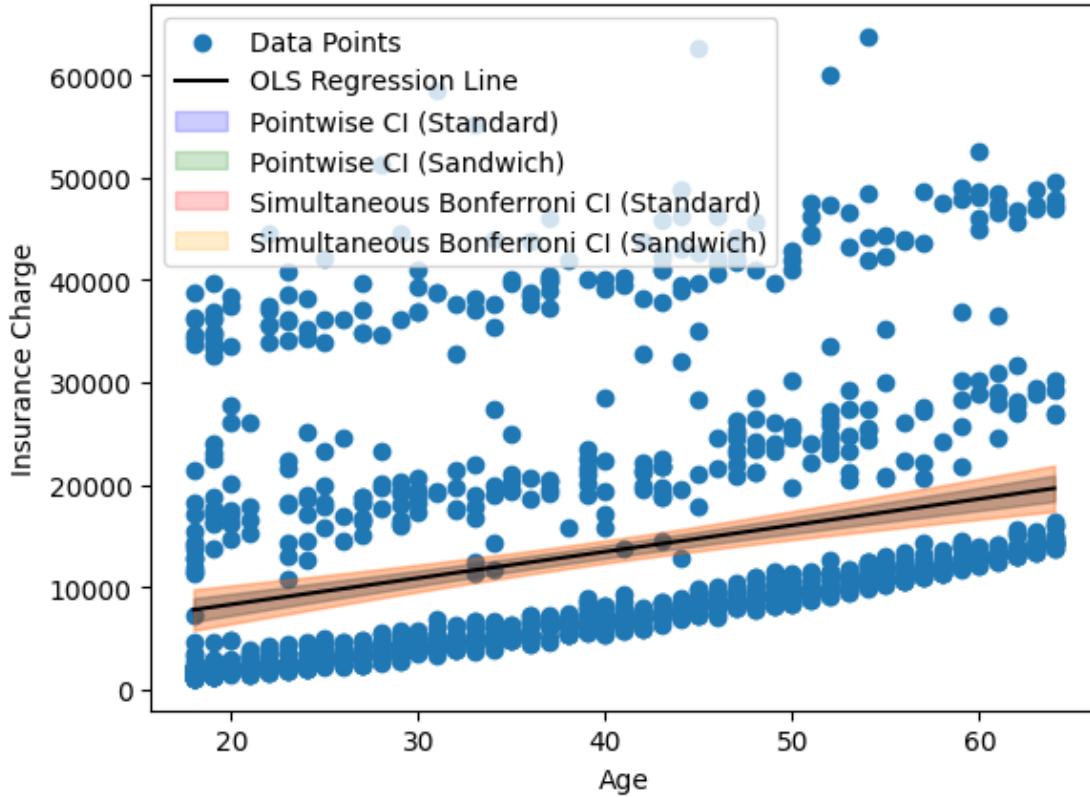
bonf_rob_u = yhat_g + z_bonf * se_rob
bonf_rob_l = yhat_g - z_bonf * se_rob

plt.scatter(x, y, label='Data Points')
plt.plot(xg, yhat_g, color='black', label='OLS Regression Line')

plt.fill_between(xg, pw_std_l, pw_std_u, color='blue', alpha=0.2, □
    ↵label='Pointwise CI (Standard)')
plt.fill_between(xg, pw_rob_l, pw_rob_u, color='green', alpha=0.2, □
    ↵label='Pointwise CI (Sandwich)')
plt.fill_between(xg, bonf_std_l, bonf_std_u, color='red', alpha=0.2, □
    ↵label='Simultaneous Bonferroni CI (Standard)')
plt.fill_between(xg, bonf_rob_l, bonf_rob_u, color='orange', alpha=0.2, □
    ↵label='Simultaneous Bonferroni CI (Sandwich)')

plt.xlabel('Age')
plt.ylabel('Insurance Charge')
plt.legend()
plt.show()

```



LLM Usage: All work was done by myself in VSCode with [GitHub Copilot integration](#). The integration “provides code suggestions, explanations, and automated implementations based on natural language prompts and existing code context,” and also offers autonomous coding and an in-IDE chat interface that is able to interact with the current codebase. Only the Copilot provided automatic inline suggestions for both LaTex and Python in `.tex` and `.ipynb` Jupyter notebook files respectively were taken into account / used.