

hw5_project

February 11, 2026

1 1

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import statsmodels.api as sm

[2]: insurance = pd.read_csv("../datasets/insurance.csv")
insurance.head()

[2]:   age      sex      bmi  children  smoker      region    charges
0    19  female  27.900        0     yes  southwest  16884.92400
1    18     male  33.770        1     no  southeast  1725.55230
2    28     male  33.000        3     no  southeast  4449.46200
3    33     male  22.705        0     no  northwest  21984.47061
4    32     male  28.880        0     no  northwest  3866.85520

[3]: X = insurance.drop(columns=['charges'])
X = pd.get_dummies(X, drop_first=True).astype(float)
y = insurance['charges']

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
print(model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          charges    R-squared:       0.751
Model:                 OLS        Adj. R-squared:   0.749
Method:                Least Squares   F-statistic:    500.8
Date:      Wed, 11 Feb 2026   Prob (F-statistic):  0.00
Time:      16:00:17            Log-Likelihood: -13548.
No. Observations:      1338        AIC:             2.711e+04
Df Residuals:          1329        BIC:             2.716e+04
Df Model:                  8
Covariance Type:    nonrobust
=====
```

====

	coef	std err	t	P> t	[0.025
0.975]					
<hr/>					
const	-1.194e+04	987.819	-12.086	0.000	-1.39e+04
-1e+04					
age	256.8564	11.899	21.587	0.000	233.514
280.199					
bmi	339.1935	28.599	11.860	0.000	283.088
395.298					
children	475.5005	137.804	3.451	0.001	205.163
745.838					
sex_male	-131.3144	332.945	-0.394	0.693	-784.470
521.842					
smoker_yes	2.385e+04	413.153	57.723	0.000	2.3e+04
2.47e+04					
region_northwest	-352.9639	476.276	-0.741	0.459	-1287.298
581.370					
region_southeast	-1035.0220	478.692	-2.162	0.031	-1974.097
-95.947					
region_southwest	-960.0510	477.933	-2.009	0.045	-1897.636
-22.466					
<hr/>					
Omnibus:	300.366	Durbin-Watson:	2.088		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	718.887		
Skew:	1.211	Prob(JB):	7.86e-157		
Kurtosis:	5.651	Cond. No.	311.		
<hr/>					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Intercept: -11940, SE 987.819, p < 0.0001 This is the predicted insurance charge for a customer age 0, bmi 0, 0 children female, non smoking, and from the north east. Realistically this doesn't mean anything, but anchors the regression line.

Age: 256.8564, SE 11.899, p < 0.0001 Holding all other predictors constant, a 1 year increase in age is attributed to a \$256.86 increase in medical insurance charge. The SE and p value are both low, indicating a well estimated and statistically significant predictor.

BMI: 339.1935, SE 28.599, p < 0.0001 Holding all other predictors constant, a 1-unit increase in BMI is associated with an average increase of \$339.19 in insurance charges. The SE is modest relative to the coefficient, suggesting reasonable precision. The p-value is extremely small, indicating strong statistical significance.

Children: 475.5005, SE 137.804, p 0.001 Holding all other predictors constant, each additional child is associated with an average increase of \$475.50 in insurance charges. The SE is somewhat larger relative to the coefficient than with age or BMI, indicating more variability in the estimate. However, the p-value (0.001) shows this predictor is still statistically significant.

Sex (Male): -131.3144, SE 332.945, p 0.693 Holding all other predictors constant, being male (relative to female) is associated with an average decrease of \$131.31 in insurance charges. However, the SE is large compared to the coefficient, and the p-value is very high (0.693). Thus, the predictor is not statistically significant.

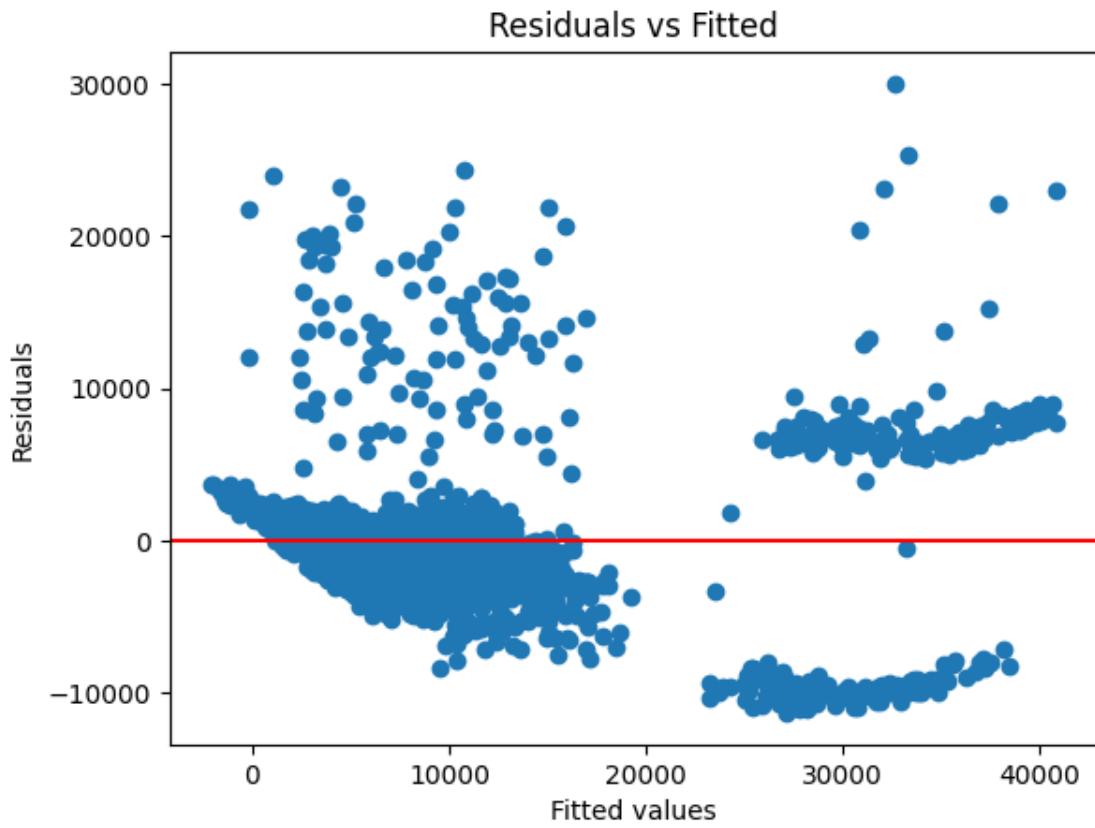
Smoker (Yes): 23850, SE 413.153, p < 0.0001 Holding all other predictors constant, being a smoker is associated with an average increase of approximately \$23,850 in insurance charges compared to non-smokers. The SE is small relative to the coefficient, and the p-value is extremely small, indicating statistical significance. This is by far the strongest predictor in the model due to its effect size and being precisely estimated.

Region (Northwest): -352.9639, SE 476.276, p 0.459 Compared to the Northeast (reference region), living in the Northwest is associated with an average decrease of \$352.96 in charges, holding other variables constant. The SE is larger than the coefficient magnitude and the p-value is high, thus the predictor is not statistically significant.

Region (Southeast): -1035.0220, SE 478.692, p 0.031 Compared to the Northeast, living in the Southeast is associated with an average decrease of \$1,035.02 in charges. The SE is moderate relative to the coefficient, and the p-value (0.031) indicates statistical significance at the 5% level.

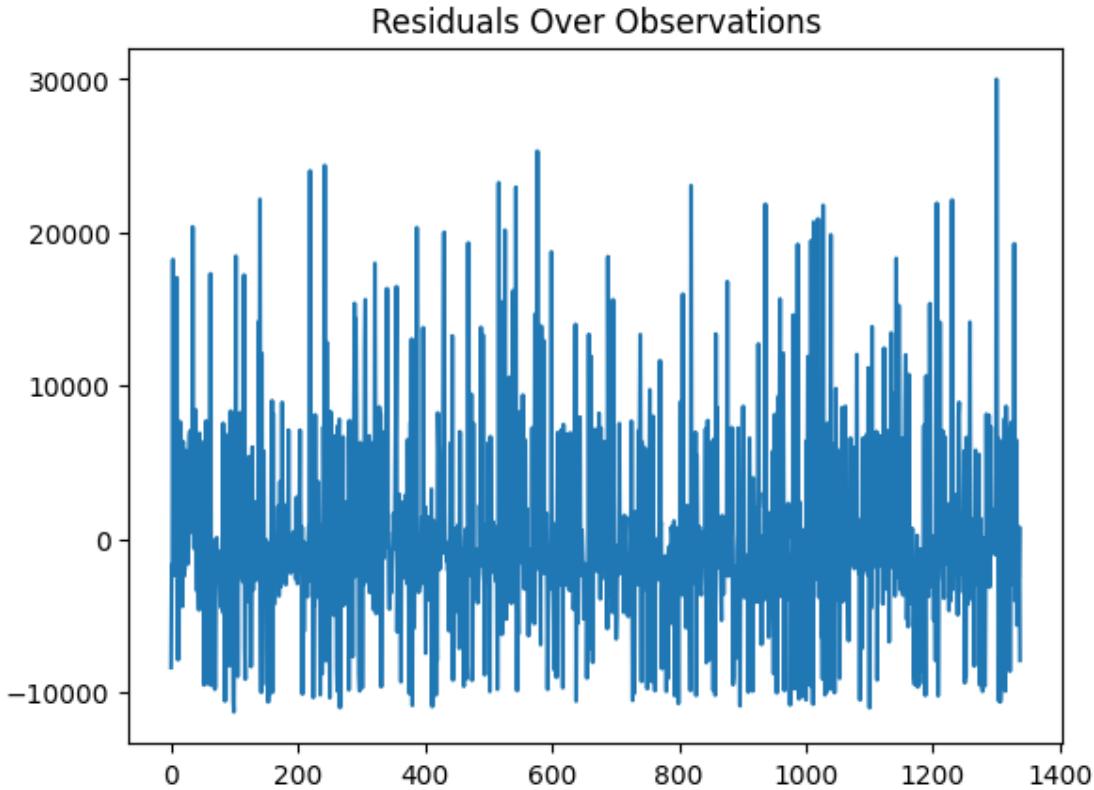
Region (Southwest): -960.0510, SE 477.933, p 0.045 Compared to the Northeast, living in the Southwest is associated with an average decrease of \$960.05 in charges. The SE is about half the magnitude of the coefficient. The p-value (0.045) indicates marginal statistical significance at the 5% level.

```
[4]: plt.scatter(model.fittedvalues, model.resid)
plt.axhline(0, color='red')
plt.xlabel("Fitted values")
plt.ylabel("Residuals")
plt.title("Residuals vs Fitted")
plt.show()
```



Residuals are clearly not random, indicating a non linear relationship in the model. This is expected as we can tell from the omnibus / skew statistics.

```
[6]: plt.plot(model.resid)
plt.title("Residuals Over Observations")
plt.show()
```



This is a good indication that there is no autocorrelation, as also indicated by the DW statistic near 2. Insurance data should not be related to time.

2 2

2.1 a

```
[8]: X = insurance.drop(columns=['charges'])
X = pd.get_dummies(X, drop_first=True).astype(float)
X = X[['age', 'bmi', 'children', 'smoker_yes']]

X['obese'] = (X['bmi'] > 30).astype(float)
X = X.drop(columns=['bmi'])

X_smoker = X[X['smoker_yes'] == 1].drop(columns=['smoker_yes'])
y_smoker = y[X['smoker_yes'] == 1]
X_smoker = sm.add_constant(X_smoker)
model_smoker = sm.OLS(y_smoker, X_smoker).fit()
print(model_smoker.summary())

X_non_smoker = X[X['smoker_yes'] == 0].drop(columns=['smoker_yes'])
```

```

y_non_smoker = y[X['smoker_yes'] == 0]
X_non_smoker = sm.add_constant(X_non_smoker)
model_non_smoker = sm.OLS(y_non_smoker, X_non_smoker).fit()
print(model_non_smoker.summary())

```

OLS Regression Results

Dep. Variable:	charges	R-squared:	0.885	
Model:	OLS	Adj. R-squared:	0.883	
Method:	Least Squares	F-statistic:	690.3	
Date:	Wed, 11 Feb 2026	Prob (F-statistic):	2.88e-126	
Time:	16:22:26	Log-Likelihood:	-2655.3	
No. Observations:	274	AIC:	5319.	
Df Residuals:	270	BIC:	5333.	
Df Model:	3			
Covariance Type:	nonrobust			
<hr/>				
	coef	std err	t	
			P> t	
			[0.025	
			0.975]	
<hr/>				
const	1.088e+04	757.111	14.377	0.000
age	272.0695	17.208	15.811	0.000
children	163.4200	206.862	0.790	0.430
obese	1.999e+04	477.353	41.874	0.000
<hr/>				
Omnibus:		187.780	Durbin-Watson:	1.796
Prob(Omnibus):		0.000	Jarque-Bera (JB):	1424.355
Skew:		2.858	Prob(JB):	5.07e-310
Kurtosis:		12.596	Cond. No.	134.
<hr/>				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

Dep. Variable:	charges	R-squared:	0.408	
Model:	OLS	Adj. R-squared:	0.407	
Method:	Least Squares	F-statistic:	243.9	
Date:	Wed, 11 Feb 2026	Prob (F-statistic):	2.55e-120	
Time:	16:22:26	Log-Likelihood:	-10485.	
No. Observations:	1064	AIC:	2.098e+04	
Df Residuals:	1060	BIC:	2.100e+04	
Df Model:	3			
Covariance Type:	nonrobust			
<hr/>				
	coef	std err	t	
			P> t	
			[0.025	
			0.975]	
<hr/>				
const	-2710.9336	449.266	-6.034	0.000
				-3592.486
				-1829.381

age	265.1057	10.108	26.228	0.000	245.272	284.939
children	580.3462	116.320	4.989	0.000	352.102	808.590
obese	135.0719	284.874	0.474	0.635	-423.909	694.053
<hr/>						
Omnibus:	708.588	Durbin-Watson:	2.058			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	5640.544			
Skew:	3.182	Prob(JB):	0.00			
Kurtosis:	12.312	Cond. No.	136.			
<hr/>						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

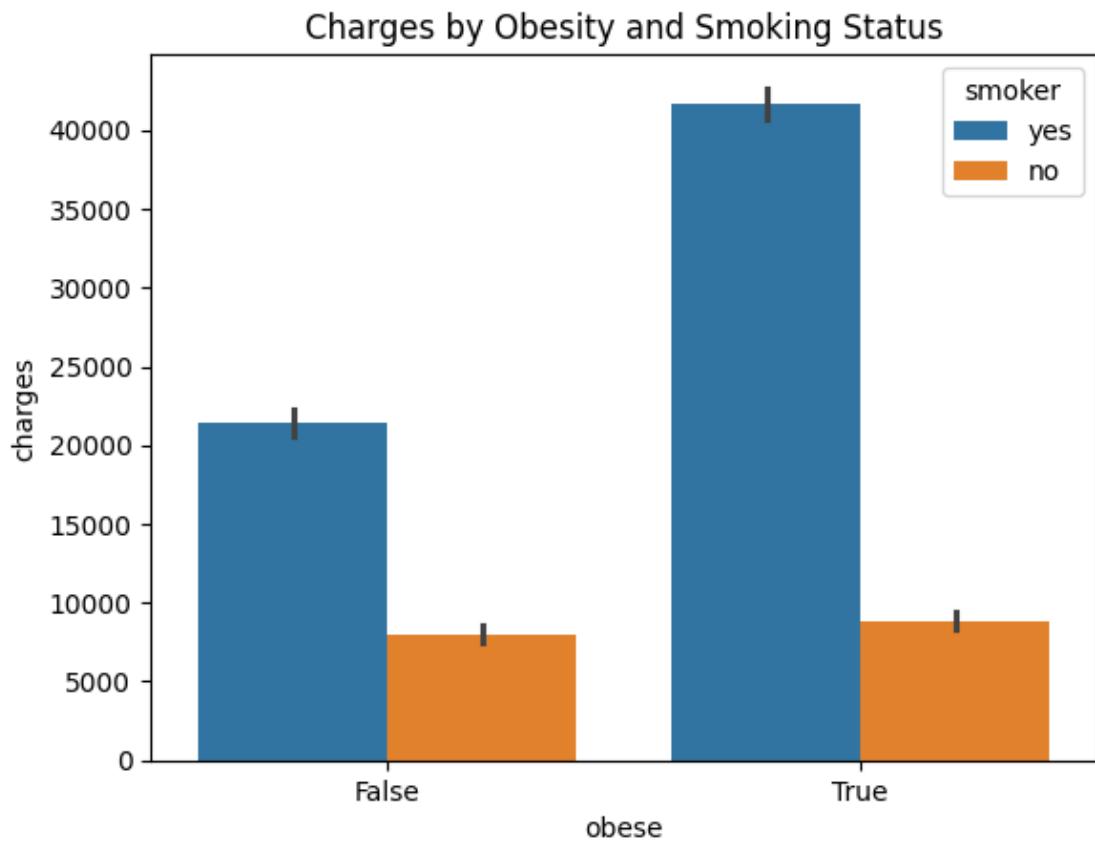
We can see that for a smoker who is also obese ($bmi > 30$), the coefficient for BMI is 19990, with an SE of 477 and a p-value < 0.0001 . However, for a non smoker and the coefficient for obese is 135 with high SE and high p-value. The difference in coefficient and statistical significance clearly indicates the presence of an interaction term that is needed.

```
[9]: import seaborn as sns

df_plot = insurance.copy()
df_plot["obese"] = (df_plot["bmi"] > 30)

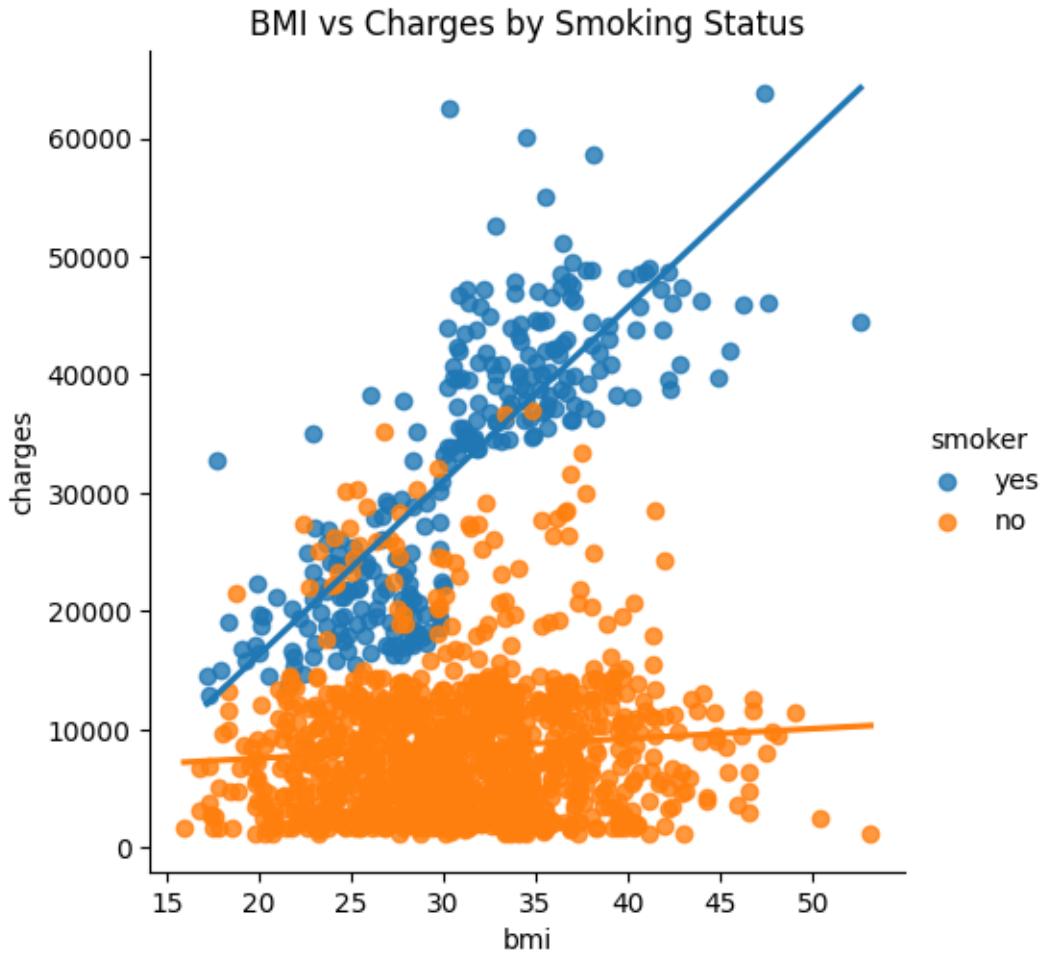
sns.barplot(
    data=df_plot,
    x="obese",
    y="charges",
    hue="smoker"
)

plt.title("Charges by Obesity and Smoking Status")
plt.show()
```



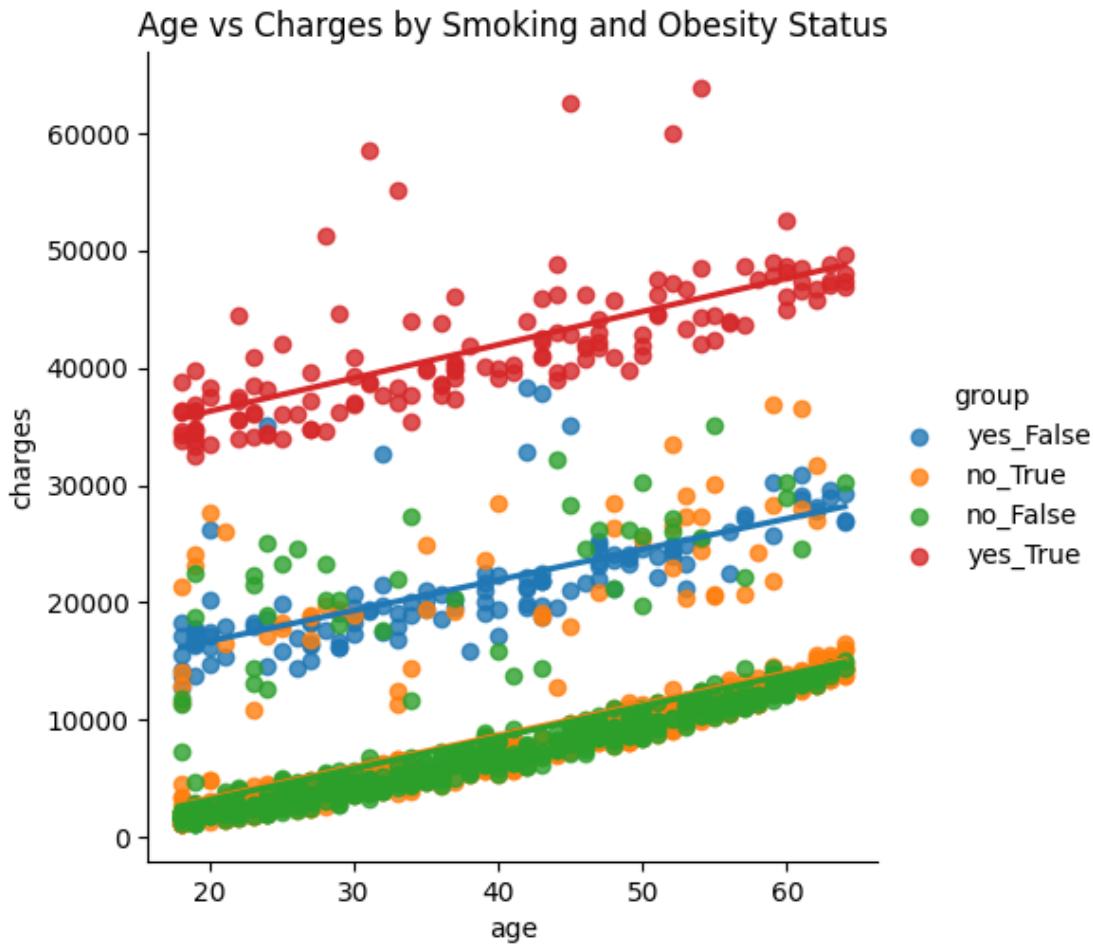
```
[10]: sns.lmplot(
    data=df_plot,
    x="bmi",
    y="charges",
    hue="smoker",
    ci=None
)

plt.title("BMI vs Charges by Smoking Status")
plt.show()
```



```
[12]: df_plot["group"] = (
    df_plot["smoker"] + "_" + df_plot["obese"].astype(str)
)

sns.lmplot(
    data=df_plot,
    x="age",
    y="charges",
    hue="group",
    ci=None
)
plt.title("Age vs Charges by Smoking and Obesity Status")
plt.show()
```



The graphs distinctly show the separation between people who smoke and are obese and their respective insurance charges. The interaction term between smoking and obesity is crucial to the model.

2.2 b

```
[13]: X = insurance.drop(columns=['charges'])
X = pd.get_dummies(X, drop_first=True).astype(float)
X = X[['age', 'bmi', 'children', 'smoker_yes']]

X['obese'] = (X['bmi'] > 30).astype(float)
X = X.drop(columns=['bmi'])

# Add interaction term
X['smoker_obese'] = X['smoker_yes'] * X['obese']
X = sm.add_constant(X)
model_interaction = sm.OLS(y, X).fit()
```

```
print(model_interaction.summary())
```

OLS Regression Results

Dep. Variable:	charges	R-squared:	0.863			
Model:	OLS	Adj. R-squared:	0.863			
Method:	Least Squares	F-statistic:	1680.			
Date:	Wed, 11 Feb 2026	Prob (F-statistic):	0.00			
Time:	16:27:18	Log-Likelihood:	-13147.			
No. Observations:	1338	AIC:	2.631e+04			
Df Residuals:	1332	BIC:	2.634e+04			
Df Model:	5					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-2670.0560	399.933	-6.676	0.000	-3454.623	-1885.489
age	266.2188	8.783	30.310	0.000	248.988	283.449
children	502.6480	101.941	4.931	0.000	302.666	702.630
smoker_yes	1.34e+04	441.692	30.327	0.000	1.25e+04	1.43e+04
obese	135.0550	276.711	0.488	0.626	-407.781	677.891
smoker_obese	1.987e+04	609.192	32.612	0.000	1.87e+04	2.11e+04
Omnibus:	885.600	Durbin-Watson:	2.060			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7358.552			
Skew:	3.167	Prob(JB):	0.00			
Kurtosis:	12.585	Cond. No.	246.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The interaction effect model is stronger, with an R^2 of 0.863 compared to the separate fitted models R^2 of 0.885 and 0.408. Although the R^2 is slightly lower, we still prefer the interaction model. The split models showed that obesity effect is massive for smokers, while negligible for non smokers. However, the models reduce sample size and prevent formal testing of a slope difference. For non-obese individuals a person who smokes (smoking coefficient) on average is predicted to have \$13400 more in insurance charges. This predictor is statistically significant. For non-smokers and obese person on average is predicted only to have a \$135 increase in charge. With a p-value of 0.626, this is not statistically significant. But, we can see that the interaction term has a coefficient of 19870, the most impactful predictor in the model, and has a low SE and extremely low p-value indicating statistical significance. While we can tell from the smaller smoking and obese coefficients that these factors alone matter, the interaction term is what affects the model the strongest. Thus, we prefer the interaction term model.

2.3 c

The interaction term specifically represents the predicted increase in insurance charges associated with obesity, which is \$19,870 larger for smokers than for non-smokers, while holding age and number of children constant.

LLM Usage: All work was done by myself in VSCode with [GitHub Copilot integration](#). The integration “provides code suggestions, explanations, and automated implementations based on natural language prompts and existing code context,” and also offers autonomous coding and an in-IDE chat interface that is able to interact with the current codebase. Only the Copilot provided automatic inline suggestions for both LaTex and Python in `.tex` and `.ipynb` Jupyter notebook files respectively were taken into account / used.